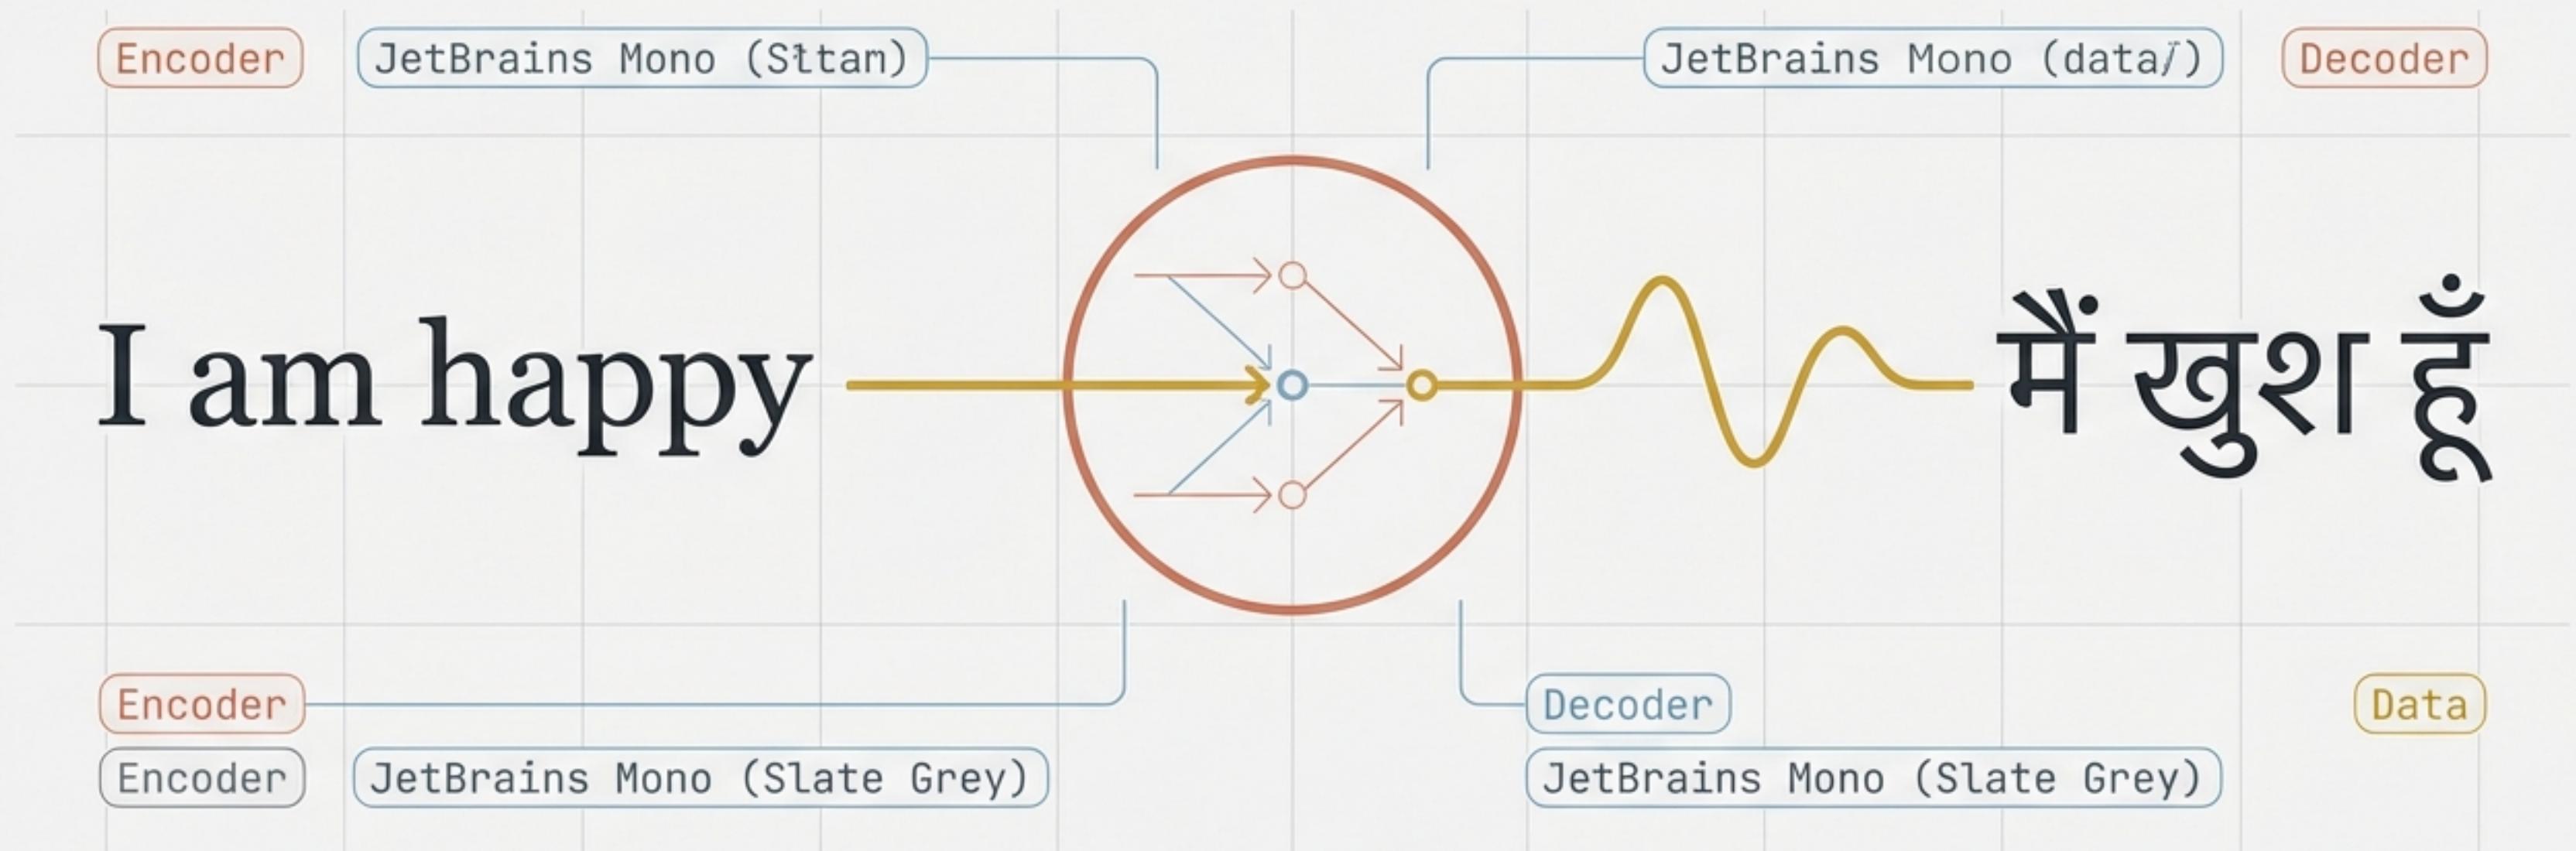


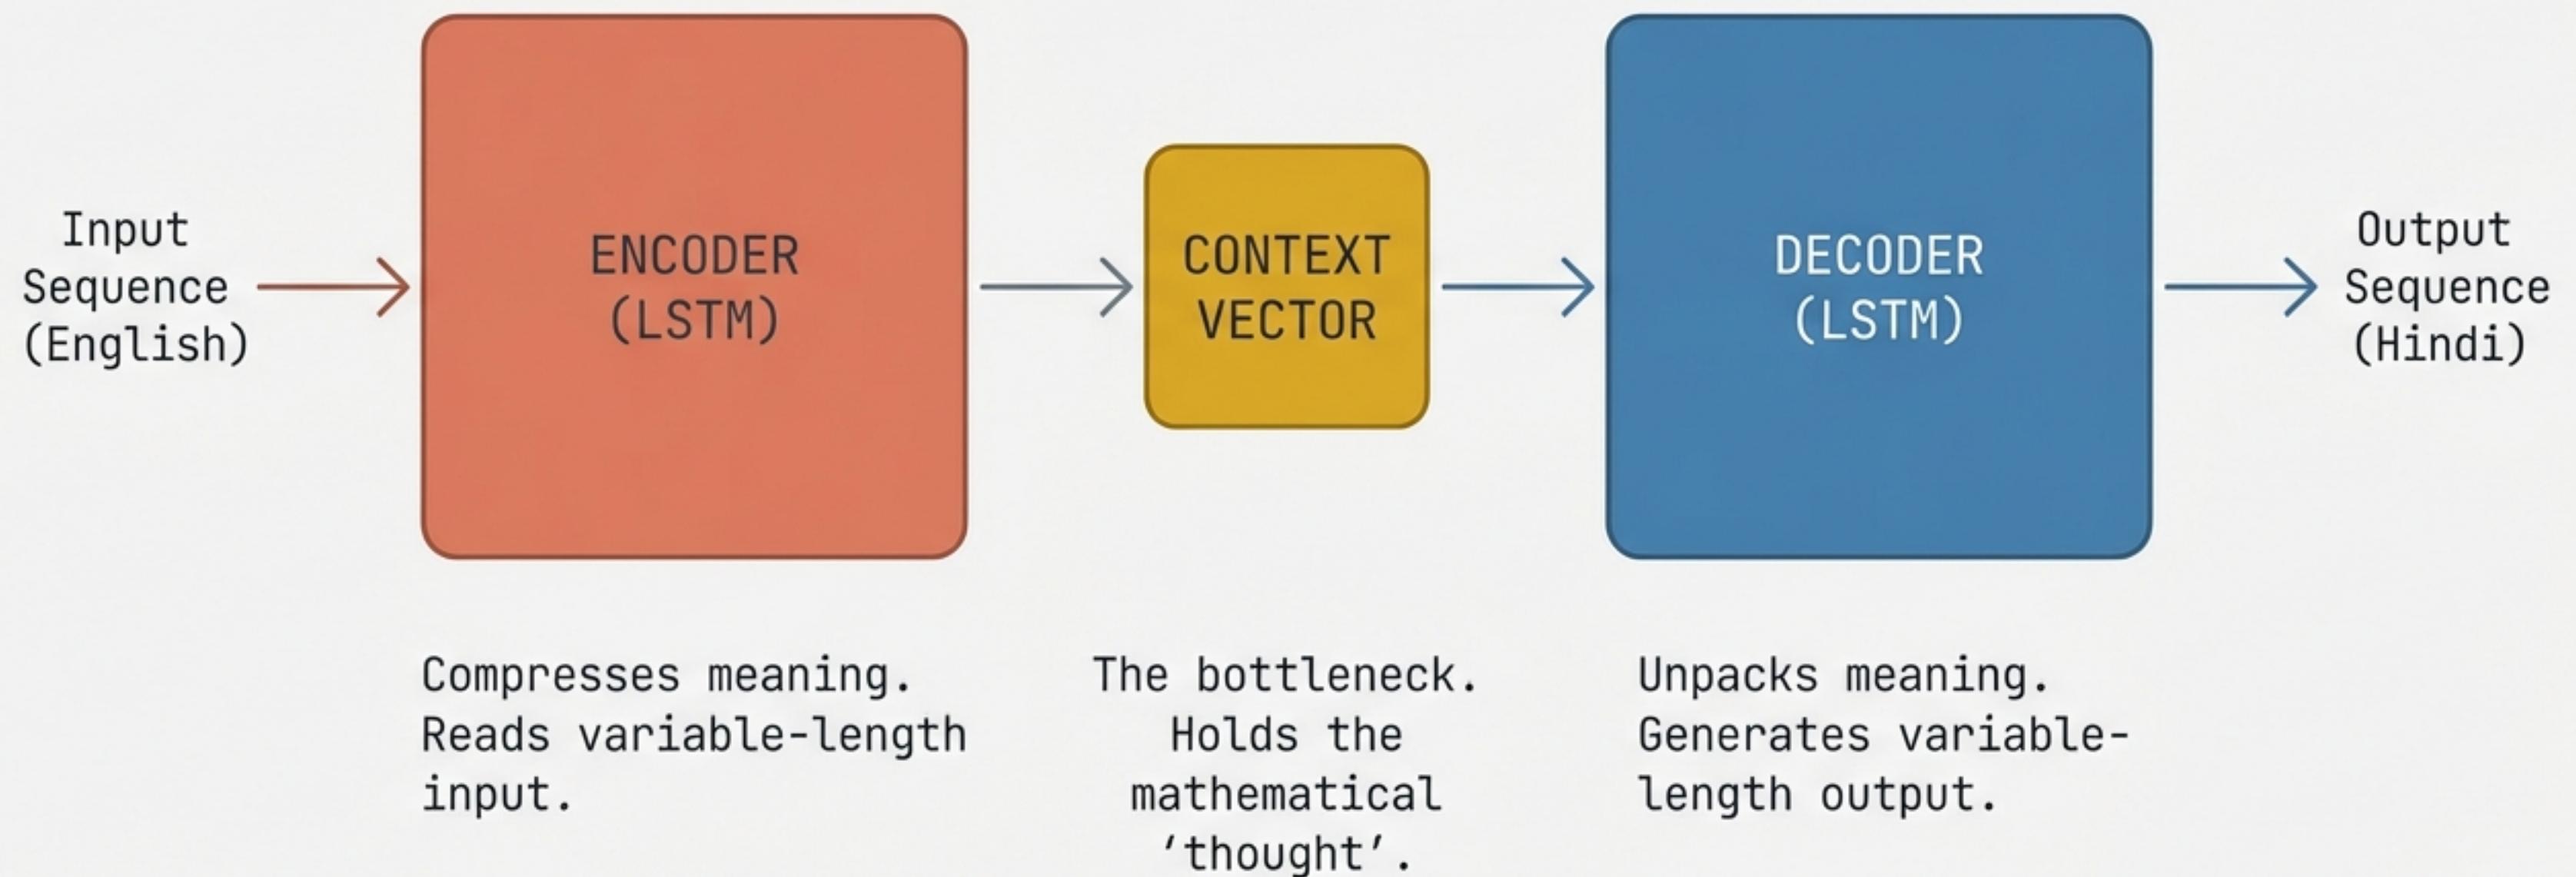
Building a Translator

A step-by-step journey inside a Seq2Seq LSTM Model.



Architecture: Sequence-to-Sequence with Attention-free LSTM · Framework: TensorFlow 2.x

The Architecture: The Map of the Journey



The Setup: Imports and Raw Data

Input: 'I am happy' | Target: 'मैं खुश हूँ'

Input: 'You are sad' | Target: 'तुम उदास हो'

```
import tensorflow as tf
import numpy as np

# The Dataset (English -> Hindi)
data = [
    ("I am happy", "मैं खुश हूँ"),
    ("You are sad", "तुम उदास हो"),
    ("She is tired", "वह थकी हुई है"),
    # ... typically millions of pairs
]
```

Vocabulary: Building the Dictionary

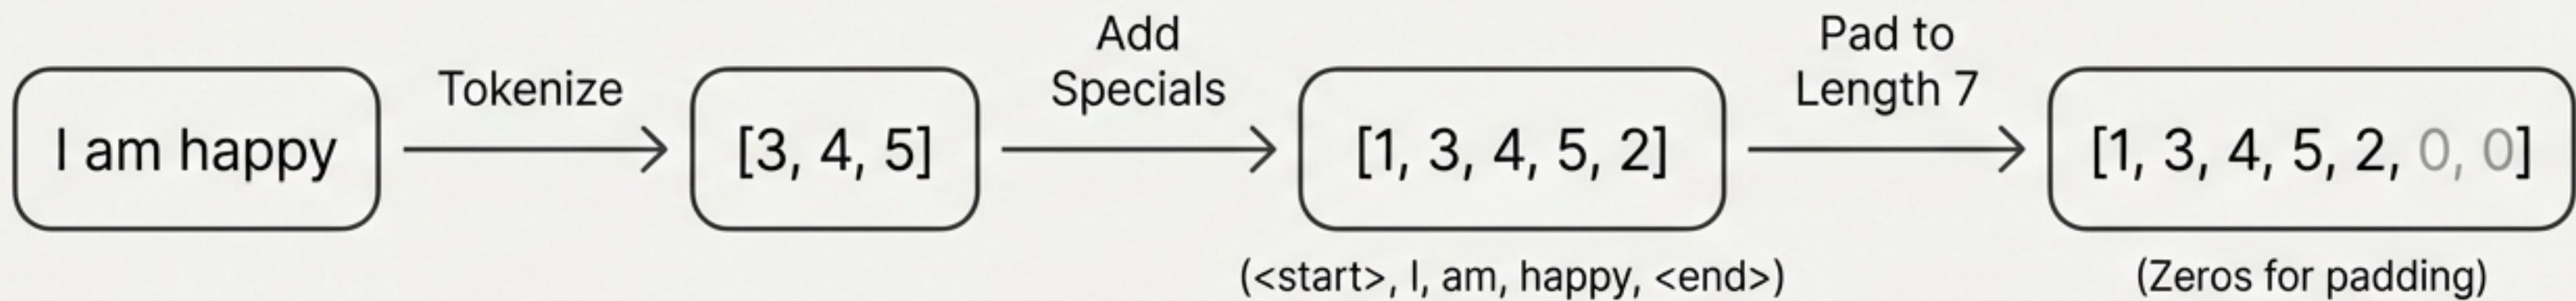
Lookup Table

Key		Value
<pad>	→	0
<start>	→	1
<end>	→	2
I	→	3
am	→	4
happy	→	5

```
def build_vocab(sentences):
    # Create mapping of words to unique IDs
    vocab = {"<pad>": 0, "<start>": 1, "<end>": 2}

    # Assign IDs to remaining words
    for sentence in sentences:
        for word in sentence.split():
            if word not in vocab:
                vocab[word] = len(vocab)
    return vocab
```

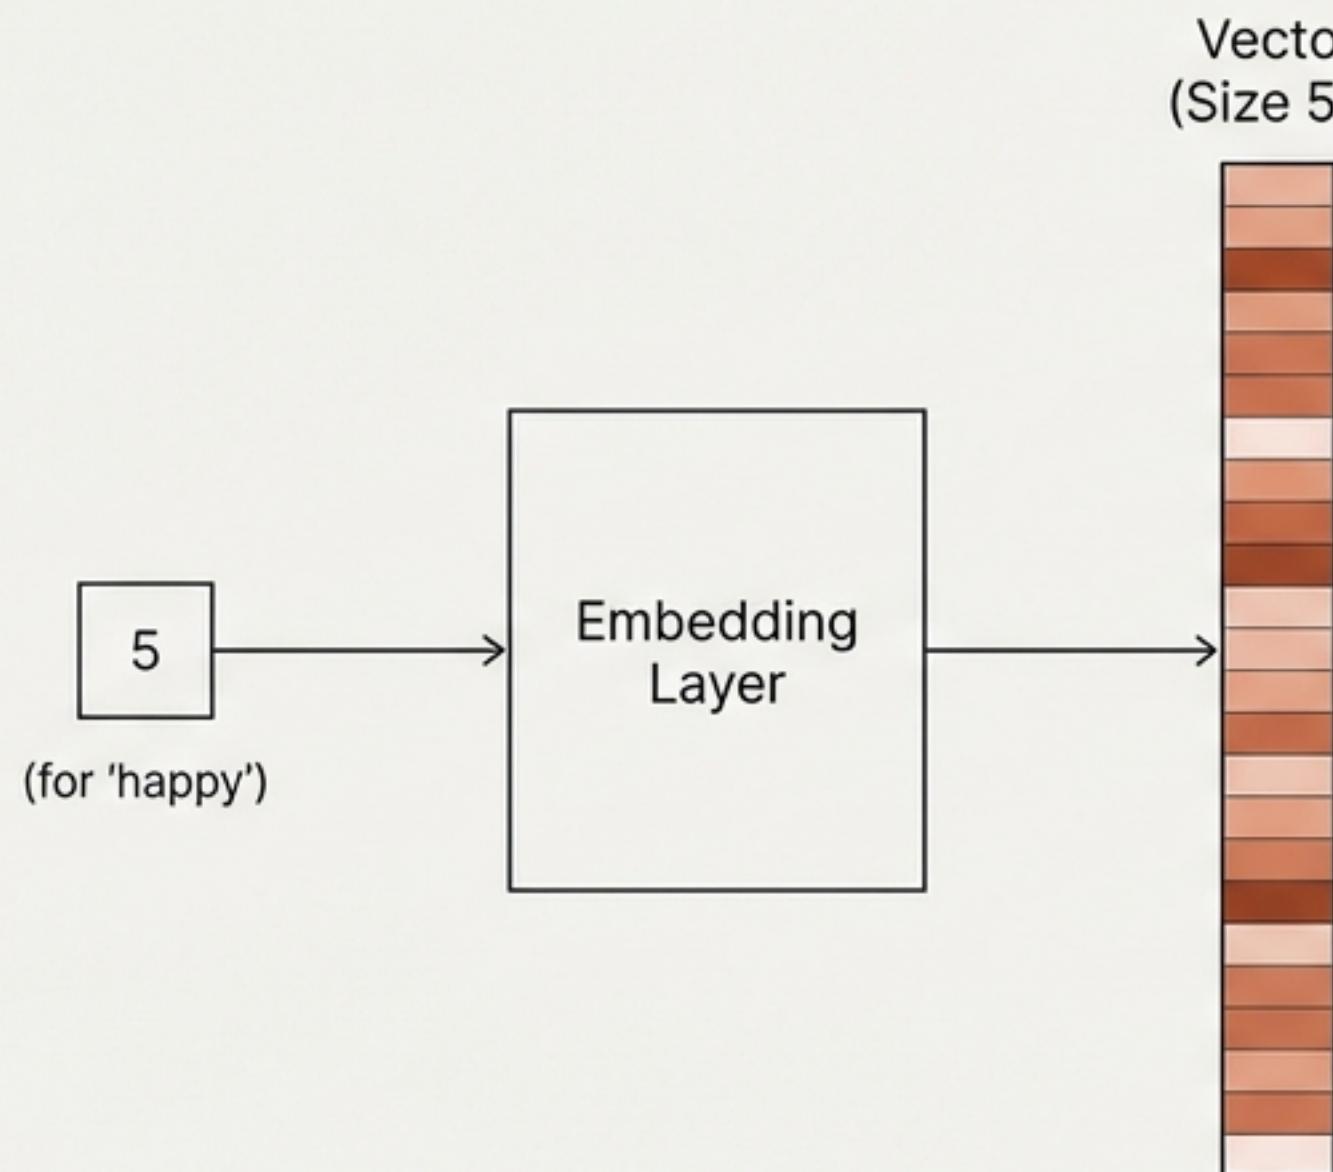
Pre-processing: Indexing and Padding



```
# Convert text to sequence of integers
src_data = [sentence_to_indices(pair[0], vocab) for pair in data]

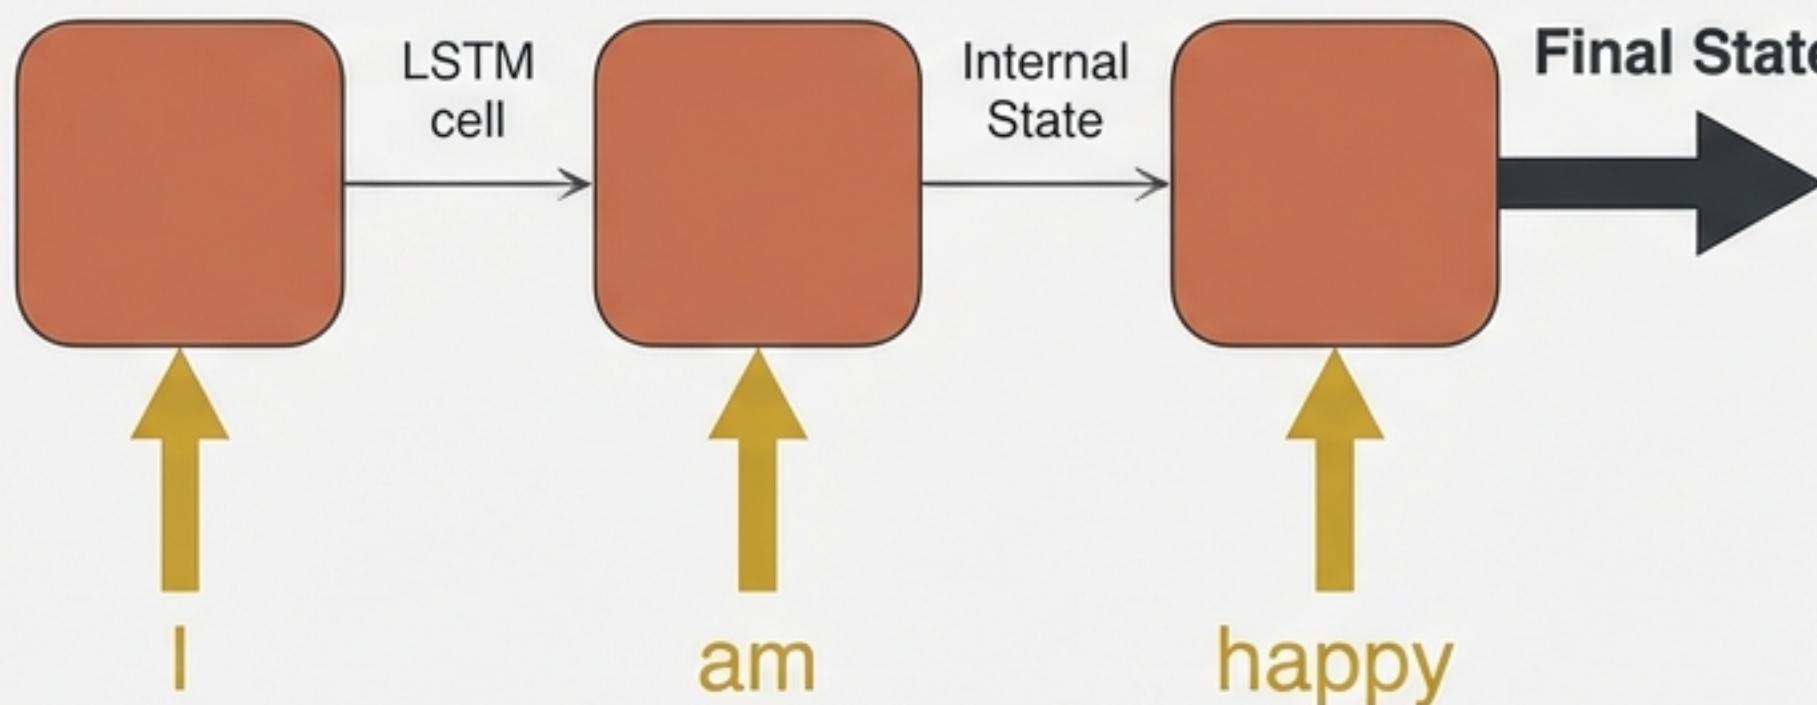
# Pad sequences to ensure equal length (post-padding)
# 0 is the <pad> token
src_padded = tf.keras.preprocessing.sequence.pad_sequences(
    src_data,
    padding='post',
    value=vocab["<pad>"]
)
```

The Encoder Part 1: The Embedding Layer



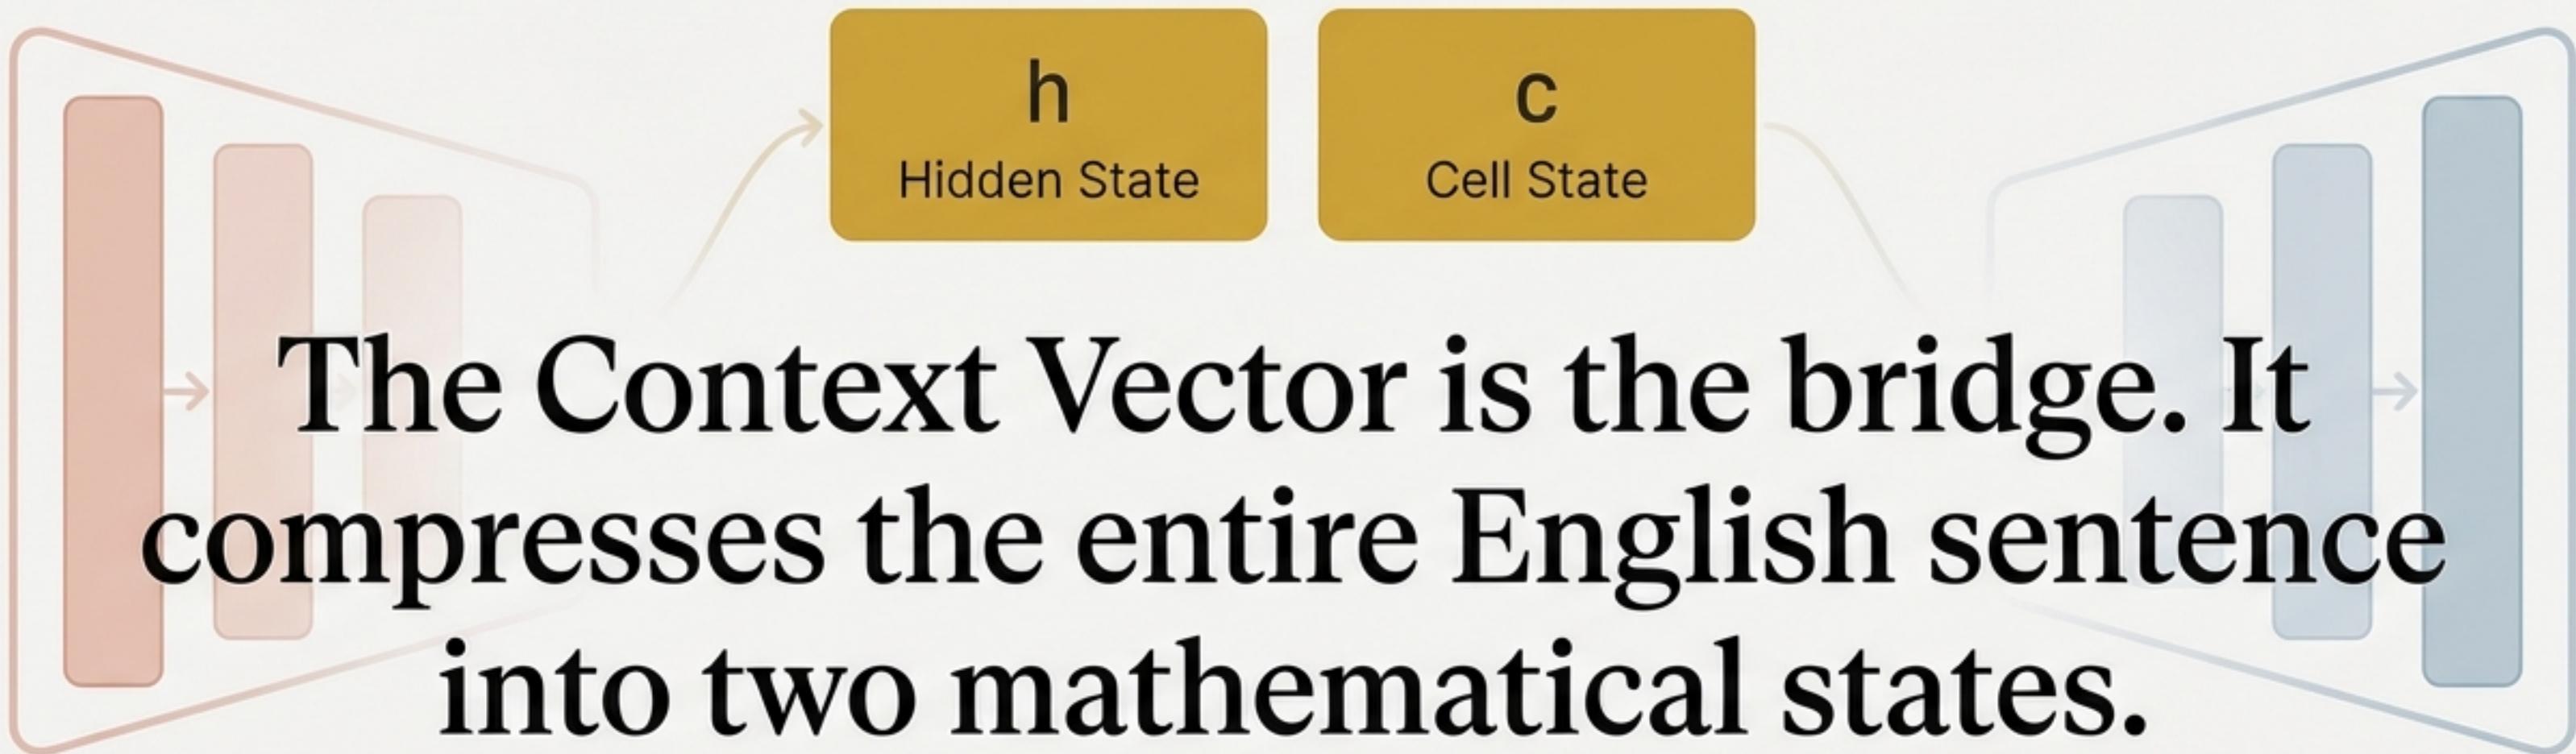
```
1  class Encoder(tf.keras.Model):
2      def __init__(self, input_size, embed_size, hidden_size):
3          super(Encoder, self).__init__()
4          # Turns indices (5) into vectors (size 50)
5          self.embedding = tf.keras.layers.Embedding(input_size, embed_size)
6
7      def call(self, x):
8          # x shape: [batch_size, seq_len]
9          embedded = self.embedding(x)
10         # embedded shape: [batch_size, seq_len, embed_size]
11         return embedded
```

The Encoder Part 2: The LSTM



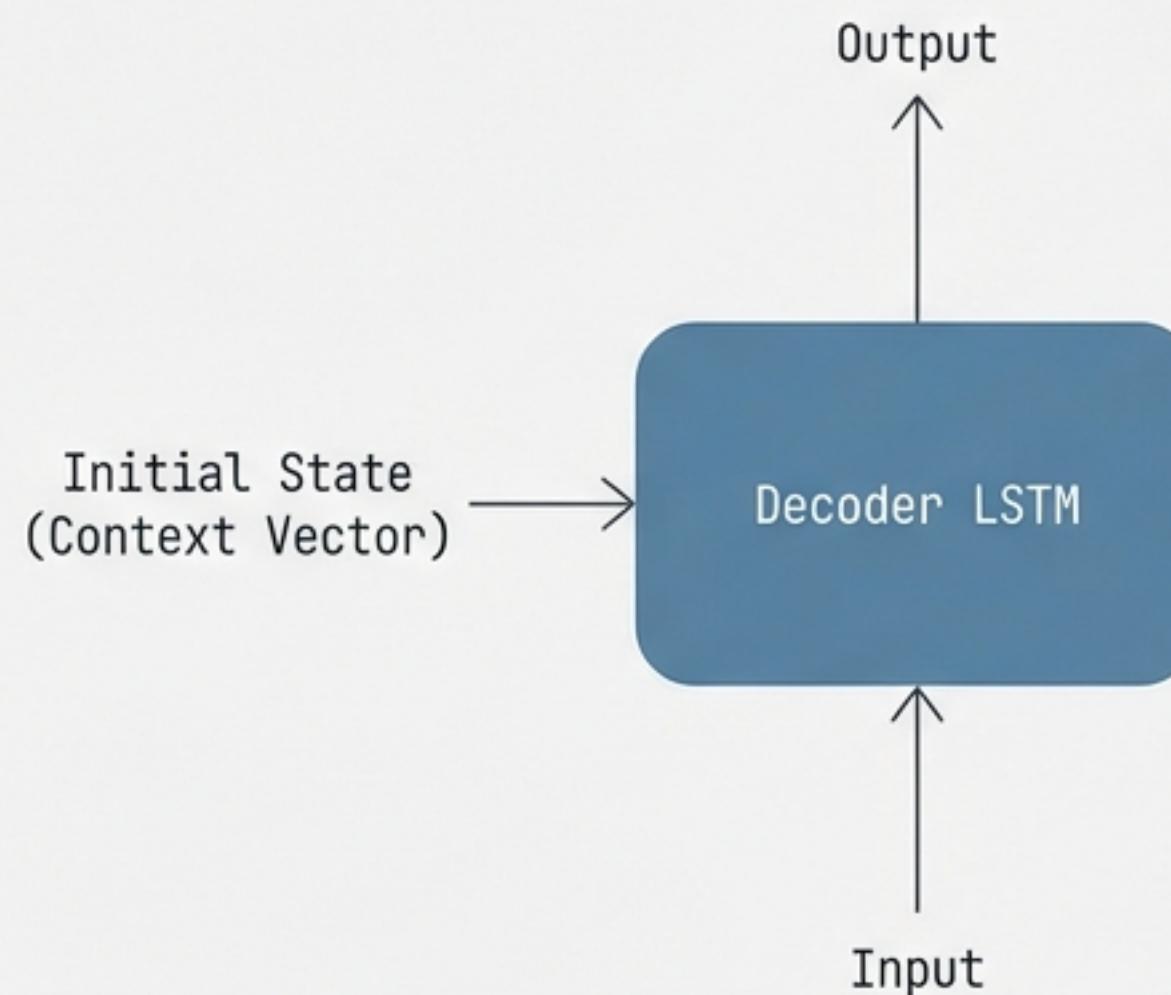
```
1 def __init__(self, ...):
2     # return_state=True is crucial!
3     self.lstm = tf.keras.layers.LSTM(hidden_size,
4         |   return_state=True)
5
6 def call(self, x):
7     embedded = self.embedding(x)
8
9     # We ignore 'output' (predictions)
10    # We keep 'hidden' and 'cell' (the memory)
11    output, hidden, cell = self.lstm(embedded)
12
13    return hidden, cell
```

The Handoff: The Context Vector



```
# From Encoder  
return hidden, cell  
  
# To Decoder  
decoder_initial_state = [hidden, cell]
```

The Decoder: Initialization

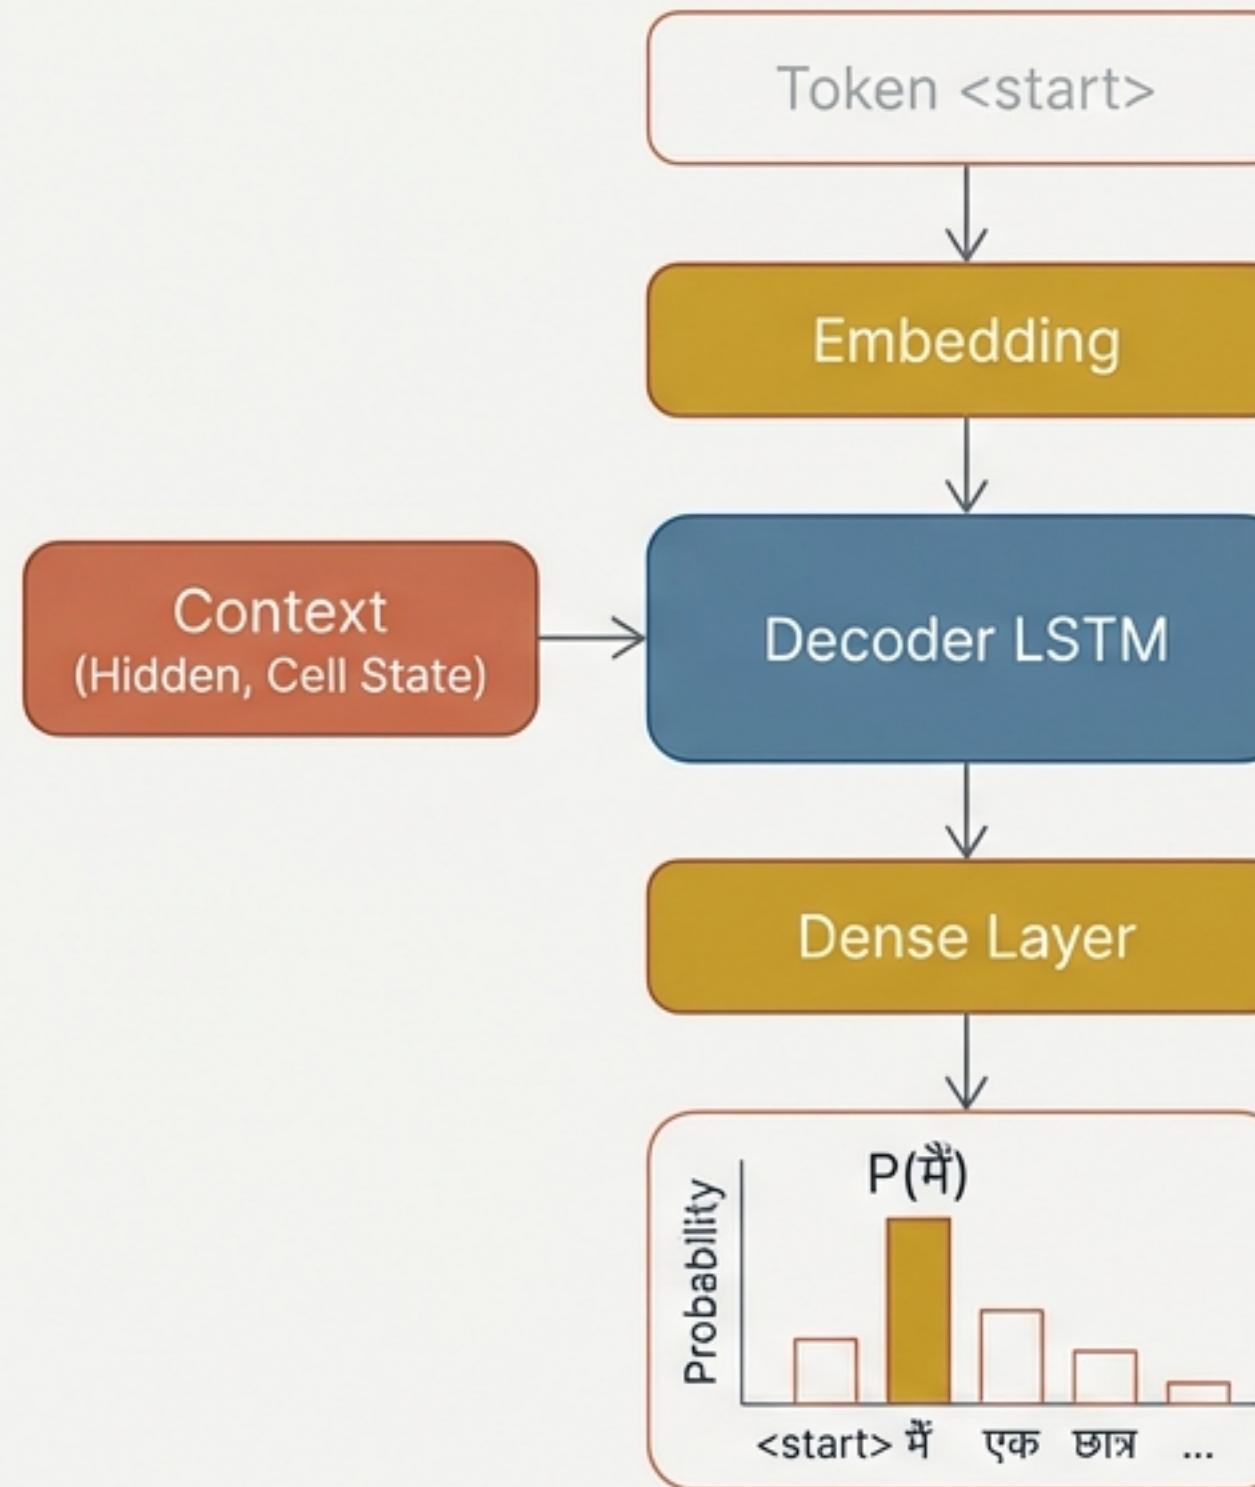


```
class Decoder(tf.keras.Model):
    def __init__(self, output_size, embed_size, hidden_size):
        super(Decoder, self).__init__()
        self.embedding = tf.keras.layers.Embedding(output_size, embed_size)

        # return_sequences=True: We need a word at every step
        self.lstm = tf.keras.layers.LSTM(hidden_size,
                                         return_sequences=True,
                                         return_state=True)

        self.fc = tf.keras.layers.Dense(output_size)
```

The Decoder: Generating Words



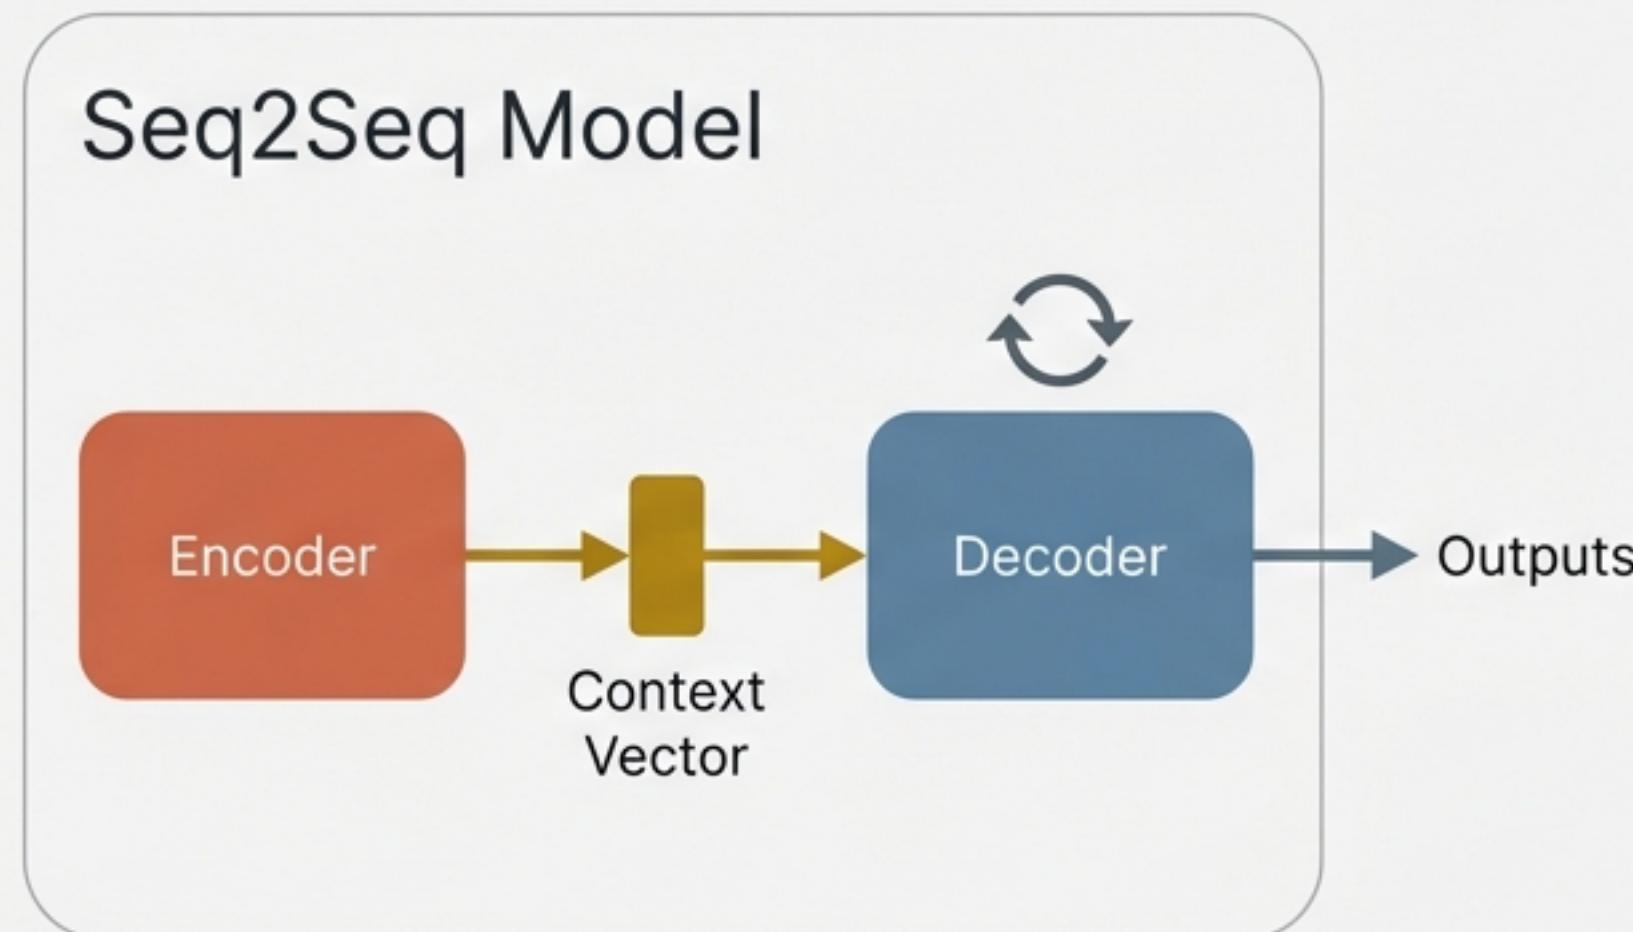
```
def call(self, x, hidden, cell):
    # x is the previous word (or <start>)
    embedded = self.embedding(x)

    # LSTM updates memory, produces output
    lstm_out, hidden, cell = self.lstm(embedded,
                                        initial_state=[hidden, cell])

    # Dense layer predicts next word probability
    output = self.fc(lstm_out)

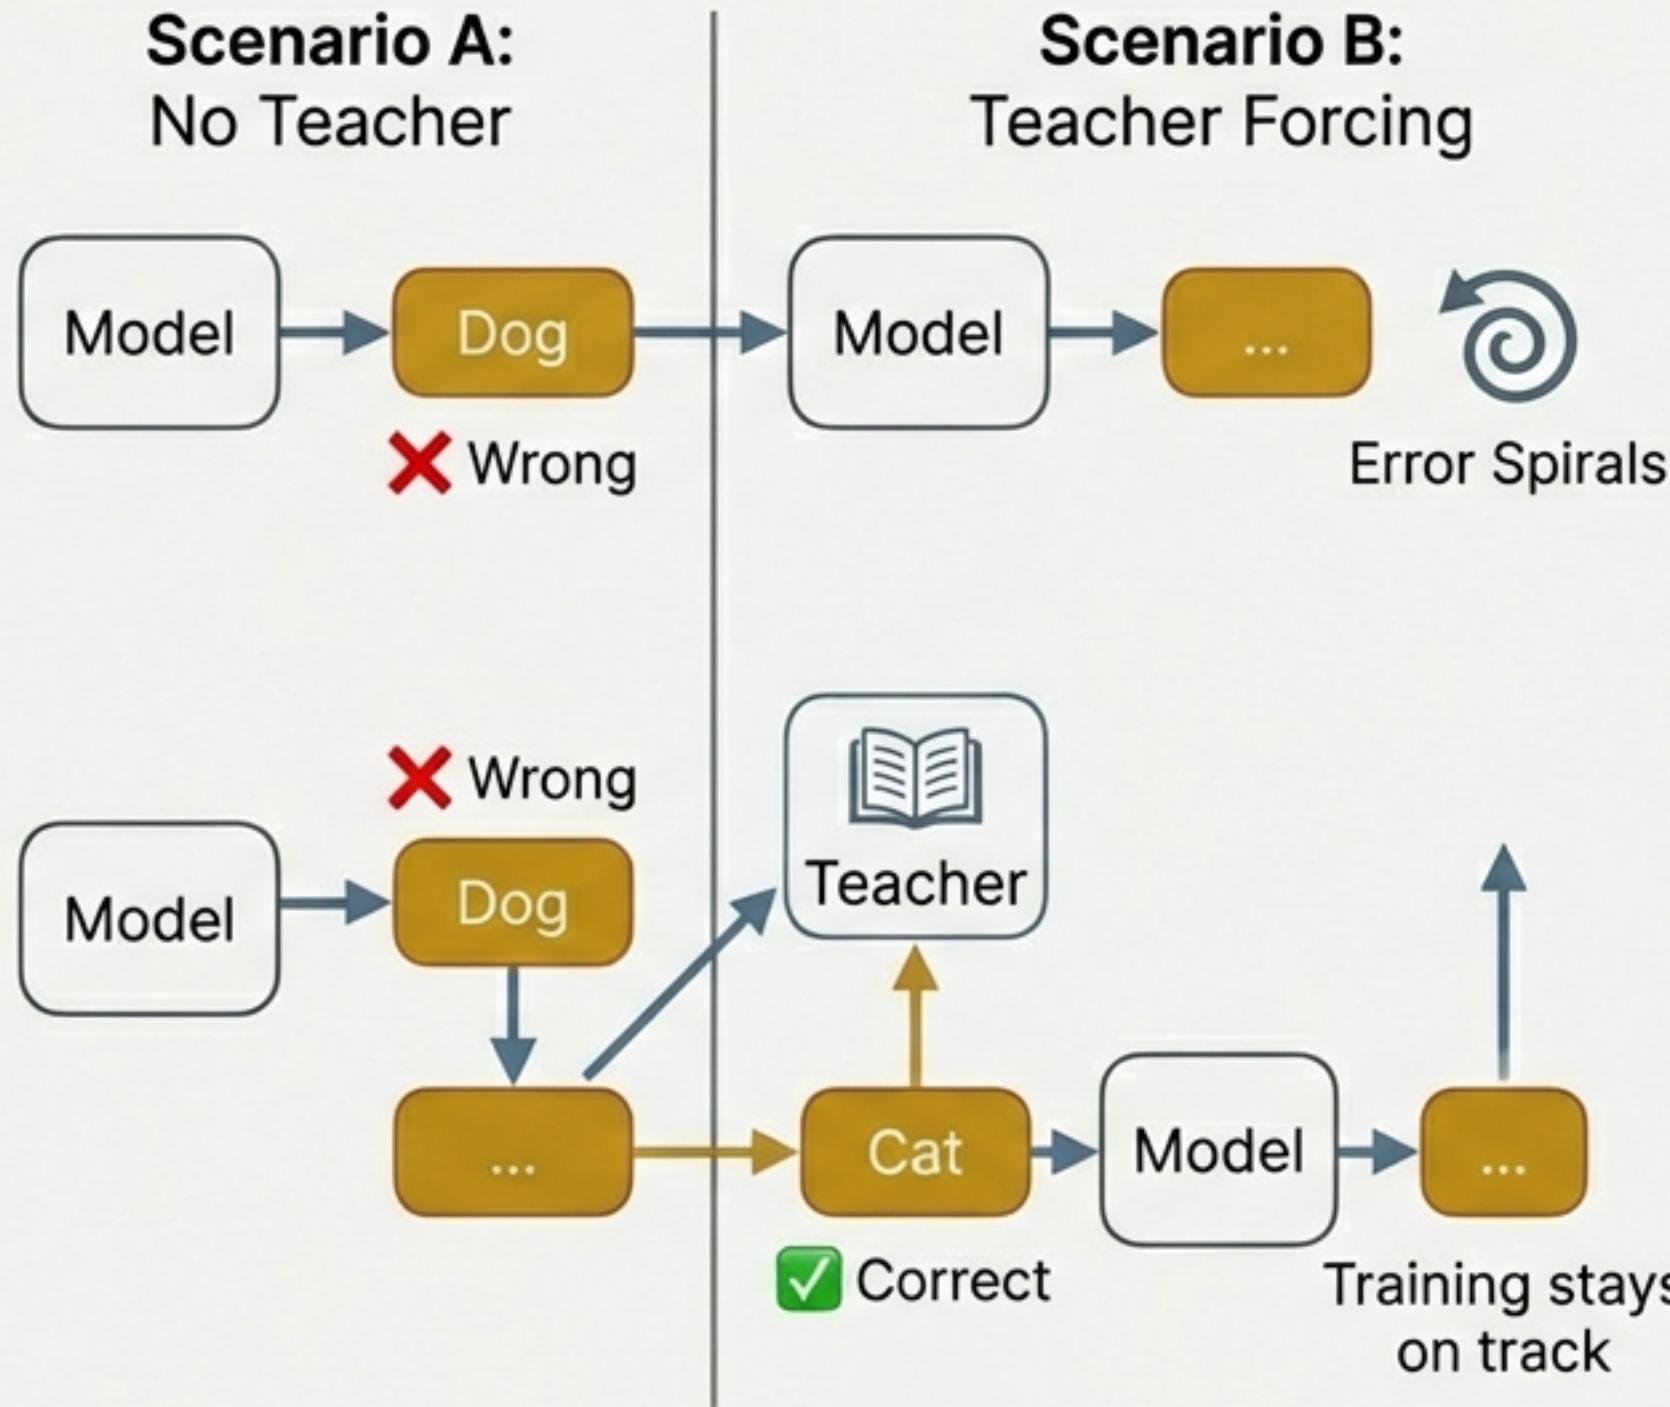
    return output, hidden, cell
```

The Wrapper: The Seq2Seq Class



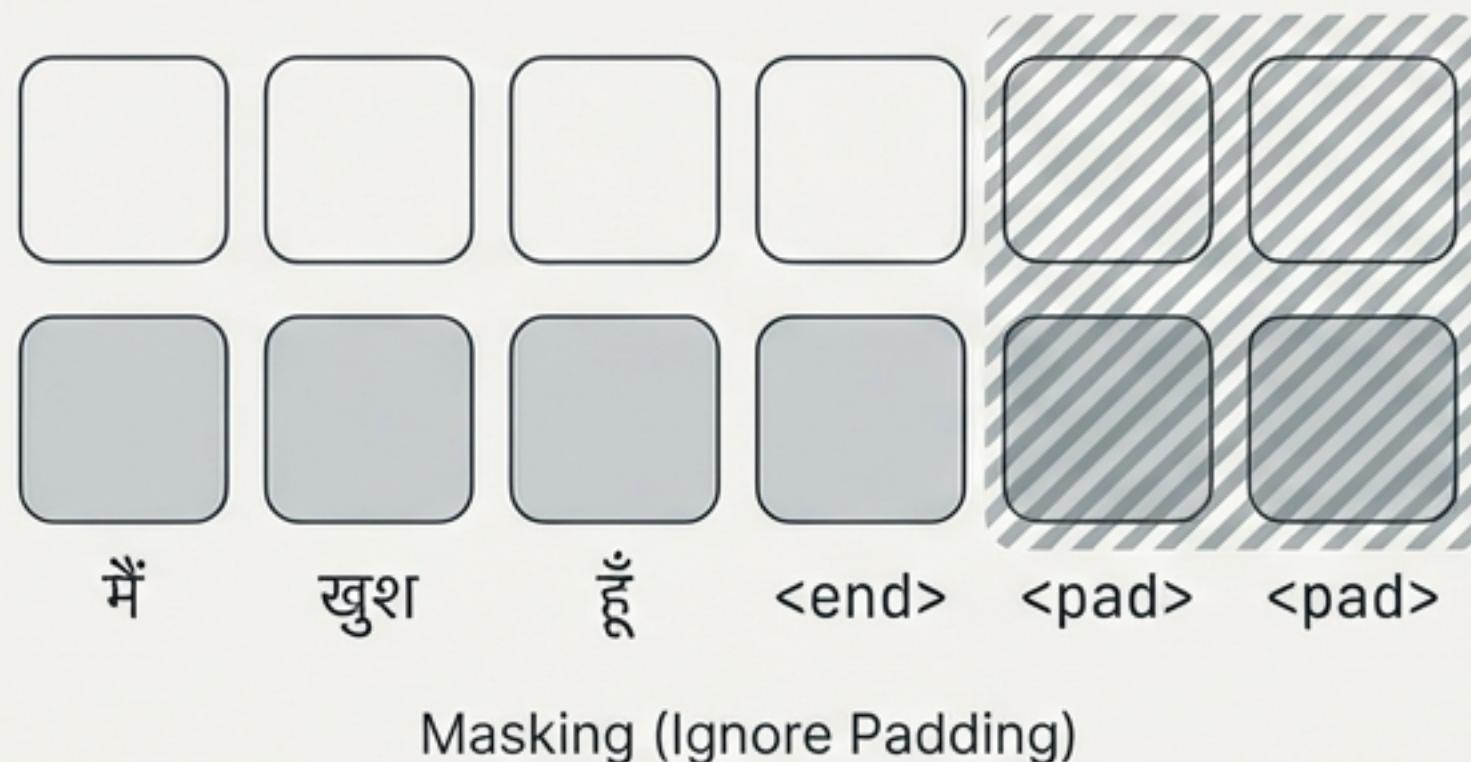
```
class Seq2Seq(tf.keras.Model):  
    def call(self, inputs):  
        src, tgt = inputs  
  
        # 1. Encode the source  
        hidden, cell = self.encoder(src)  
  
        # 2. Decode loop  
        outputs = []  
        input_token = tgt[:, 0:1] # Start with <start>  
  
        for t in range(1, tgt_len):  
            output, hidden, cell =  
                self.decoder(input_token, hidden, cell)  
            outputs.append(output)  
            # ... update input_token for next step ...  
  
        return tf.concat(outputs, axis=1)
```

Training Strategy: Teacher Forcing



```
if training:  
    # Teacher Forcing: Feed the ACTUAL  
    # next word from dataset  
    teacher_force = random.random() <  
        teacher_forcing_ratio  
    input_token = tgt[:, t:t+1] if  
        teacher_force else  
        predicted_token  
  
else:  
    # Inference: Must use model's own  
    # prediction  
    input_token = predicted_token
```

The Loss Function: Grading the Exam



```
loss_fn = tf.keras.losses.SparseCategoricalCrossentropy(  
    from_logits=True, reduction='none')  
  
# Create mask to ignore padding tokens (0)  
# We don't want to penalize the model for padding errors  
mask = tf.cast(target_labels != 0, tf.float32)  
  
loss = loss_fn(target_labels, outputs)  
loss = loss * mask # Apply mask to zero-out padding loss
```

Inference: Translation in Action

Input: "I am happy"

Step 1: Predicted "मैं"

Step 2: Predicted "खुश"

Step 3: Predicted "हूँ"

Step 4: Predicted "<end>" -> STOP

Result: "मैं खुश हूँ"

```
def translate(sentence):
    # Encode inputs...
    hidden, cell = model.encoder(src_tensor)
    input_token = [start_token]

    for step in range(max_len):
        output, hidden, cell = model.decoder(
            input_token, hidden, cell)
        predicted_token = tf.argmax(output,
                                     axis=-1)

        if predicted_token == end_token:
            break
        input_token = predicted_token
```

Concepts Mastered

	Tokenization: Converting text to indices. JetBrains Mono, Slate Grey		Encoder-Decoder: The Seq2Seq architecture. JetBrains Mono, Slate Grey
	Padding: Handling variable sequence lengths. JetBrains Mono, Slate Grey		Context Vector: The information bottleneck. JetBrains Mono, Slate Grey
	Embedding: Dense vector representation of meaning. JetBrains Mono, Slate Grey		Teacher Forcing: Stabilizing training. JetBrains Mono, Slate Grey

This Seq2Seq + LSTM architecture laid the groundwork for modern NLP. While Attention mechanisms and Transformers (BERT, GPT) have since evolved, they stand on the shoulders of this logic.